

Generative Models - DRE7053

Lecture 4

NORA Summer School 2024

Rogelio A Mancisidor
Assistant Professor
Department of Data Science and Analytics
BI Norwegian Business School

June 10-14, 2024

1 Diffusion Models

- Denoising Diffusion Models
- Forward Diffusion Process
- Backward Diffusion Process

2 The Objective Function

- Deriving the ELBO

1 Diffusion Models


- Denoising Diffusion Models
- Forward Diffusion Process
- Backward Diffusion Process

2 The Objective Function

- Deriving the ELBO

Forward and Backward steps

- Denoising diffusion models has two steps:
 - 1 The forward diffusion process
 - 2 The backward process, which learns to generate the data



The Joint Distribution of the Diffusion Process

- At each step in the forward diffusion process we have

$$q(\mathbf{x}_t|\mathbf{x}_{t-1}) = \mathcal{N}(\sqrt{1 - \beta_t}\mathbf{x}_{t-1}, \beta_t\mathbf{I})$$

- Hence, the joint probability of the entire process from \mathbf{x}_0 (original sample) to \mathbf{x}_t is

$$q(\mathbf{x}_{1:T}|\mathbf{x}_0) = \prod_{t=1}^T q(\mathbf{x}_t|\mathbf{x}_{t-1}) \quad (1)$$

where the variance β_t increases with t

- Note that each step depends only on its previous state (a Markovian process!)
- Which is similar to the inference model in hierarchical VAEs

The Joint Distribution of the Diffusion Process


- At each step, we use the same sampling mechanism as in VAEs:

- Hence we can sample at each t from

$$\mathbf{x}_t \sim \mathcal{N}(\sqrt{\bar{\alpha}_t}\mathbf{x}_0, (1 - \bar{\alpha}_t)\mathbf{I}) \quad (2)$$

- β_t schedule is designed such that $\bar{\alpha}_T \rightarrow 0$, so $q(\mathbf{x}_T|\mathbf{x}_0) \approx \mathcal{N}(\mathbf{0}, \mathbf{1})$

Forward Diffusion Process



- Note that

- In the forward step we do not train any neural network
- The distribution $q(x_T|x_0)$ is the same as the prior distribution in VAEs

- The forward diffusion process is an iterative procedure that transforms the input data \mathbf{x}_0
- We add *noise* to \mathbf{x}_0 such that at time T

$$\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, I)$$

- That is, the mean and variance of \mathbf{x}_0 moves gradually towards $\mathbf{0}$ and $\mathbf{1}$
- The unknown distribution of \mathbf{x}_0 should be more complex than $\mathcal{N}(\mathbf{0}, I)$

Backward Diffusion Process - I

- We want to *undo* the forward diffusion process
- Ideally, we want to estimate $q(\mathbf{x}_{t-1}|\mathbf{x}_t)$ recursively
- However,

$$q(\mathbf{x}_{t-1}|\mathbf{x}_t) = \frac{q(\mathbf{x}_t|\mathbf{x}_{t-1})q(\mathbf{x}_{t-1})}{q(\mathbf{x}_t)}$$
$$\propto q(\mathbf{x}_t|\mathbf{x}_{t-1}) \int q(\mathbf{x}_{t-1}|\mathbf{x}_0)q(\mathbf{x}_0)d\mathbf{x}_0$$

it is intractable (for similar reasons as in VAEs)

Backward Diffusion Process - II

- We approximate $q(\mathbf{x}_{t-1}|\mathbf{x}_t)$ directly and create an iterative noise-reduction process $p(\mathbf{x}_{t-1}|\mathbf{x}_t) \approx q(\mathbf{x}_{t-1}|\mathbf{x}_t)$
- The approximation holds if β_t is small in the forward diffusion step
- We learn the parameters of $p(\mathbf{x}_{t-1}|\mathbf{x}_t)$, using neural networks (just as we did in VAEs), i.e.

$$p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t) = \mathcal{N}(\boldsymbol{\mu} = f_{\theta}(\mathbf{x}_t, t), \boldsymbol{\sigma}^2 = f_{\theta}(\mathbf{x}_t, t))$$

- $f_{\theta}(\cdot)$ is commonly a U-net

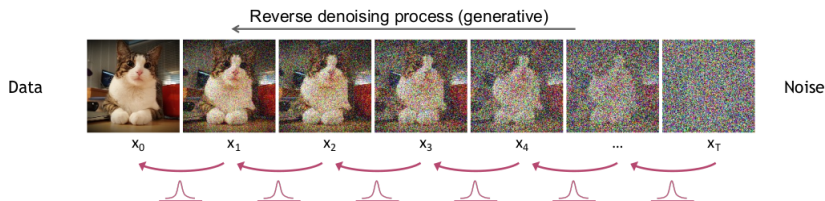
Backward Diffusion Process - III

- The joint probability of the reverse process is

$$p_{\theta}(\mathbf{x}_{0:T}) = p(\mathbf{x}_T) \prod_{t=1}^T p_{\theta}(\mathbf{x}_{t-1} | \mathbf{x}_t) \quad (3)$$

where $p(\mathbf{x}_T) \sim \mathcal{N}(\mathbf{0}, I)$

- Note that this is similar to the generative process in hierarchical VAEs



1 Diffusion Models

- Denoising Diffusion Models
- Forward Diffusion Process
- Backward Diffusion Process

2 The Objective Function

- Deriving the ELBO

Deriving the (last!) ELBO

- We start by defining an objective function that maximizes the likelihood of \mathbf{x}_0 assign by the model p_θ
- Specifically, we optimize

$$\log p(\mathbf{x}_0) \geq \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_T) \prod_{t=1}^T p(\mathbf{x}_{t-1}|\mathbf{x}_t)}{\prod_{t=1}^T q(\mathbf{x}_t|\mathbf{x}_{t-1})} \right] \quad (4)$$

which circumvents the marginal distributions that are intractable.

- Note that the expectation of a function of a subset of variables is given by taking the expectation only over the subset!

$$\mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)}[f(\mathbf{x}_{a:b})] = \mathbb{E}_{q(\mathbf{x}_{a:b}|\mathbf{x}_0)}[f(\mathbf{x}_{a:b})]$$